

DL for Case Based Reasoning - Reproduction

Berend Jansen, Tom Lotze, Stan Lochtenberg, Cees Kaandorp

FACT in AI, Group 4

January 31, 2020

Outline

1 Introduction

- Transparency in AI
- Prototype based reasoning

2 Methodology & Experiments

3 Results


- Accuracy
- Prototypes

4 Conclusion

5 Discussion

Transparency in black box models

- ▶ Decision making is inherently nontransparent in neural networks
- ▶ Visualizing model reasoning can give useful insights



(a) Husky classified as wolf



(b) Explanation

Figure 1: Raw data and explanation of model prediction [4]

Types of interpretability [2]

- ▶ Post hoc
 - ▶ LIME
 - ▶ SHAP
- ▶ Example-based
 - ▶ Counterfactual examples
 - ▶ Adversarial examples
 - ▶ Influential Instances
 - ▶ Prototypes

Prototype based reasoning

- ▶ Classification based on similarity to prototypes in latent space
- ▶ Visualize prototypes to explain model

Prototype

Typical representation of the data, showing general characteristics.

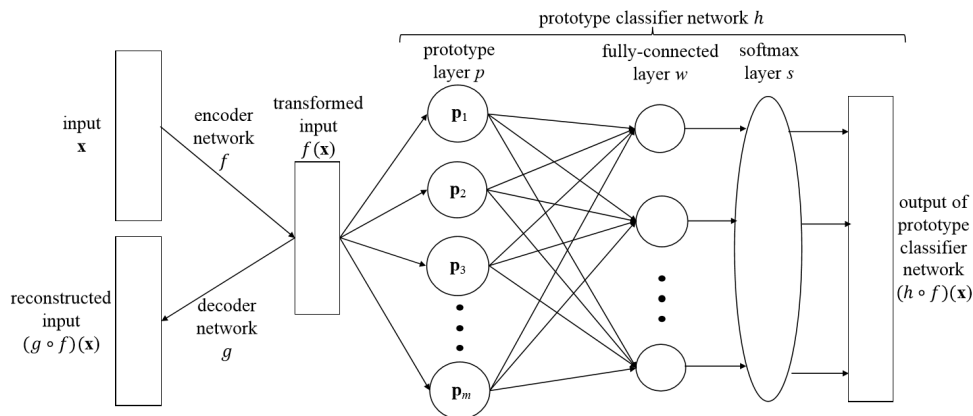


Figure 2: Prototype classification model [1]

Deep Learning for Case-Based Reasoning through Prototypes: A Neural Network that Explains Its Predictions

Oscar Li^{*1}, Hao Liu^{*3}, Chaofan Chen¹, Cynthia Rudin^{1,2}

¹Department of Computer Science, Duke University, Durham, NC, USA 27708

²Department of Electrical and Computer Engineering, Duke University, Durham, NC, USA 27708

³Kuang Yaming Honors School, Nanjing University, Nanjing, China, 210000

runliang.li@duke.edu, 141242059@smail.nju.edu.cn, {cfchen, cynthia}@cs.duke.edu

- ▶ 3 datasets:
 - ▶ MNIST handwritten digits
 - ▶ Car angles
 - ▶ Fashion MNIST

Model: Cost function

$$L = \lambda_{class} * E + \lambda_{ae} * R + \lambda_1 * R1 + \lambda_2 * R2 \quad (1)$$

- ▶ E : Cross Entropy loss
- ▶ R : Reconstruction loss
- ▶ $R1$: Push prototypes to have meaningful decodings in pixel space
- ▶ $R2$: Cluster training examples around prototypes in latent space

Methodology

- ▶ Conversion of code base from TensorFlow to PyTorch
 - ▶ Prototype model (PM) and Linear Model (LM)
- ▶ 4 different datasets (all 60,000 images)
 - ▶ Standard MNIST digits
 - ▶ CIFAR-10
 - ▶ MNIST colored (natural) background¹
 - ▶ MNIST gray (natural) background
- ▶ Grid search for optimal λ in colored & gray MNIST




Figure 3: Lena & Examples of colored MNIST dataset

¹Inspired by <https://github.com/wouterbulten/deeplearning-resources>



Results: Accuracy on different datasets

Model on dataset	Accuracy	Nr. epochs
LM on MNIST color	0.956	40
LM on MNIST rgb2gray	0.980	20
Li et al. [1] on standard MNIST	0.995	1500
PM on standard MNIST	0.991	20
PM on MNIST color	0.947	20
PM on MNIST rgb2gray	0.977	25
PM on CIFAR-10	0.643	30



Table 1: Performance of the best model on different datasets. Accuracy is computed on hold-out test set. PM refers to the Prototype Model, LM refers to the Linear Model

Results: Prototype examples



standard MNIST



CIFAR-10



MNIST color



MNIST gray





Figure 4: Prototype examples per dataset

Lambda grid search

- ▶ Effect on accuracy marginal
- ▶ Effect on prototypes significant




Figure 5: Prototypes learned in grid search:

Left: $\lambda_{class} = 20$, $\lambda_{ae} = 10$, $\lambda_1 = 1$, $\lambda_2 = 10$;

Middle: $\lambda_{class} = 20$, $\lambda_{ae} = 10$, $\lambda_1 = 1$, $\lambda_2 = 20$;

Right: $\lambda_{class} = 20$, $\lambda_{ae} = 1$, $\lambda_1 = 10$, $\lambda_2 = 1$

Conclusion

- ▶ More complicated backgrounds can still produce meaningful prototypes
- ▶ Colored images lead to problems
- ▶ Little to no intra-class variation required for meaningful prototypes
- ▶ Badge: Results Replicated
 - ▶ "The main results of the paper have been obtained in a subsequent study by a person or team other than the authors, using, in part, artifacts provided by the authors" [3]



- ▶ FACT Li et al. [1]
 - ▶ Discrepancies between code and original paper
 - ▶ Not for all datasets code available
- ▶ Limitations of our work
 - ▶ Model cannot handle complex datasets
 - ▶ Follow up paper by original authors

Questions?

- [1] Li, O., Liu, H., Chen, C., and Rudin, C. (2018). Deep learning for case-based reasoning through prototypes: A neural network that explains its predictions. In *Thirty-Second AAAI Conference on Artificial Intelligence*.
- [2] Molnar, C. (2019). *Interpretable Machine Learning*.
<https://christophm.github.io/interpretable-ml-book/>.
- [3] of Computing Machinery, A. (2018). *Artifact review and badging*.
- [4] Ribeiro, M. T., Singh, S., and Guestrin, C. (2016). " why should i trust you?" explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*, pages 1135–1144.